

▼ 1.Importing necessary dependencies

```
!pip install seaborn==0.11
```



Collecting seaborn==0.11

Downloading seaborn-0.11.0-py3-none-any.whl (283 kB)

283.1/283.1 kB 4.8 MB/s eta 0:00:00

```
Requirement already satisfied: numpy>=1.15 in /usr/local/lib/python3.10/dist-packages (from seaborn==0.11) (1.23
Requirement already satisfied: scipy>=1.0 in /usr/local/lib/python3.10/dist-packages (from seaborn==0.11) (1.11.
Requirement already satisfied: pandas>=0.23 in /usr/local/lib/python3.10/dist-packages (from seaborn==0.11) (1.5
Requirement already satisfied: matplotlib>=2.2 in /usr/local/lib/python3.10/dist-packages (from seaborn==0.11) (
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=2.2
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=2.2->se
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=2.
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=2.
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=2.2-
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=2.2->s
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=2.2
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib>
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.23->seabc
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->m
Installing collected packages: seaborn
Attempting uninstall: seaborn
Found existing installation: seaborn 0.12.2
Uninstalling seaborn-0.12.2:
Successfully uninstalled seaborn-0.12.2
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This
lida 0.0.10 requires fastapi, which is not installed.
lida 0.0.10 requires kaleido, which is not installed.
lida 0.0.10 requires python-multipart, which is not installed.
lida 0.0.10 requires uvicorn, which is not installed.
Successfully installed seaborn-0.11.0
```

```
# for data reading and data manipulation
import numpy as np
import pandas as pd
import statistics as st

# for data visualization
import matplotlib.pyplot as plt
import seaborn as sns

# for model creation and model evaluation
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression

# reading data from a .csv file to a Pandas DataFrame
df = pd.read_csv('Customer_Conversion_Prediction.csv')
pd.set_option('display.max_columns',None)
df.head()
```

```

      age      job  marital  education_qual  call_type  day  mon  dur  num_calls  prev_outcome  y
0    58  management   married      tertiary   unknown    5  may  261         1      unknown    no
df.shape

(45211, 11)

df.columns

Index(['age', 'job', 'marital', 'education_qual', 'call_type', 'day', 'mon',
      'dur', 'num_calls', 'prev_outcome', 'y'],
      dtype='object')

print(f'► The DataFrame(df) contains {df.shape[0]} rows and {df.shape[1]} columns.')

► The DataFrame(df) contains 45211 rows and 11 columns.

```

▼ 2.Checking and Handling Missing Values

Checking for missing values

```

# looking for the amount of null data in the pandas dataframe
df.isnull().sum()

```

```

age          0
job          0
marital      0
education_qual  0
call_type    0
day          0
mon          0
dur          0
num_calls    0
prev_outcome 0
y            0
dtype: int64

```

▼ 3.Data Type Correction

```
df.dtypes
```

```

age          int64
job          object
marital      object
education_qual  object
call_type    object
day          int64
mon          object
dur          int64
num_calls    int64
prev_outcome  object
y            object
dtype: object

```

▼ 4.Analyze Categorical Columns

```
# for column "job"
```

```
df['job'].value_counts()
```

```
blue-collar    9732
management    9458
technician     7597
admin          5171
services       4154
retired        2264
self-employed  1579
entrepreneur   1487
unemployed     1303
housemaid      1240
student         938
unknown        288
Name: job, dtype: int64
```

```
df['job'] = df['job'].replace({'admin.': 'admin'})
```

```
df['job'].value_counts()
```

```
blue-collar    9732
management    9458
technician     7597
admin          5171
services       4154
retired        2264
self-employed  1579
entrepreneur   1487
unemployed     1303
housemaid      1240
student         938
unknown        288
Name: job, dtype: int64
```

```
# for column "marital"
```

```
df['marital'].value_counts()
```

```
married    27214
single     12790
divorced    5207
Name: marital, dtype: int64
```

```
# for column "education_qual"
```

```
df['education_qual'].value_counts()
```

```
secondary    23202
tertiary     13301
primary       6851
unknown       1857
Name: education_qual, dtype: int64
```

```
# for column "call_type"
```

```
df['call_type'].value_counts()
```

```
cellular    29285
unknown     13020
telephone    2906
Name: call_type, dtype: int64
```

```
# for column "mon"
```

```
df['mon'].value_counts()
```

```

may      13766
jul       6895
aug       6247
jun       5341
nov       3970
apr       2932
feb       2649
jan       1403
oct        738
sep        579
mar        477
dec        214
Name: mon, dtype: int64

```

```
df['mon'] = df['mon'].replace({'jul':'july','jun':'june','aug':'auguest','nov':'november','apr':'april','feb':'februa
```

```
df['mon'].value_counts()
```

```

may      13766
july      6895
auguest   6247
june      5341
november   3970
april      2932
february   2649
january    1403
october     738
september   579
march       477
december    214
Name: mon, dtype: int64

```

```
# for column "prev_outcome"
```

```
df['prev_outcome'].value_counts()
```

```

unknown    36959
failure     4901
other       1840
success     1511
Name: prev_outcome, dtype: int64

```

```
df = df.drop_duplicates()
```

```
df.shape
```

```
(45205, 11)
```

▼ 5.Data Visualization.

```
df.columns
```

```

Index(['age', 'job', 'marital', 'education_qual', 'call_type', 'day', 'mon',
       'dur', 'num_calls', 'prev_outcome', 'y'],
      dtype='object')

```

▼ 5.1.Distribution of Future Plot

```
#plotting count plots for all the categorical columns
sns.set_theme(style='darkgrid',palette='pastel')

plt.figure(figsize=(20,25))
plt.subplot(431)
sns.countplot(df['job'],order=df['job'].value_counts().index[::-1])
plt.xticks(rotation=15)
plt.xlabel('Job')
plt.ylabel('Count')
plt.title('Conversion Prediction of Customers according to job')

plt.subplot(432)
sns.countplot(df['marital'],order=df['marital'].value_counts().index[::-1])
plt.xticks(rotation=15)
plt.xlabel('Marital')
plt.ylabel('Count')
plt.title('Conversion Prediction of Customers according to marital')

plt.subplot(433)
sns.countplot(df['education_qual'],order=df['education_qual'].value_counts().index[::-1])
plt.xticks(rotation=15)
plt.xlabel('Education Qualification')
plt.ylabel('Count')
plt.title('Conversion Prediction of Customers according to Education Qualification')

plt.subplot(434)
sns.countplot(df['call_type'],order=df['call_type'].value_counts().index[::-1])
plt.xticks(rotation=15)
plt.xlabel('Call Type')
plt.ylabel('Count')
plt.title('Conversion Prediction of Customers according to Call Type')

plt.subplot(435)
sns.countplot(df['mon'],order=df['mon'].value_counts().index[::-1])
plt.xticks(rotation=15)
plt.xlabel('Month')
plt.ylabel('Count')
plt.title('Conversion Prediction of Customers according to Month')

plt.subplot(436)
sns.countplot(df['prev_outcome'],order=df['prev_outcome'].value_counts().index[::-1])
plt.xticks(rotation=15)
plt.xlabel('Prev outcome')
plt.ylabel('Count')
plt.title('Conversion Prediction of Customers according to Prev outcome')


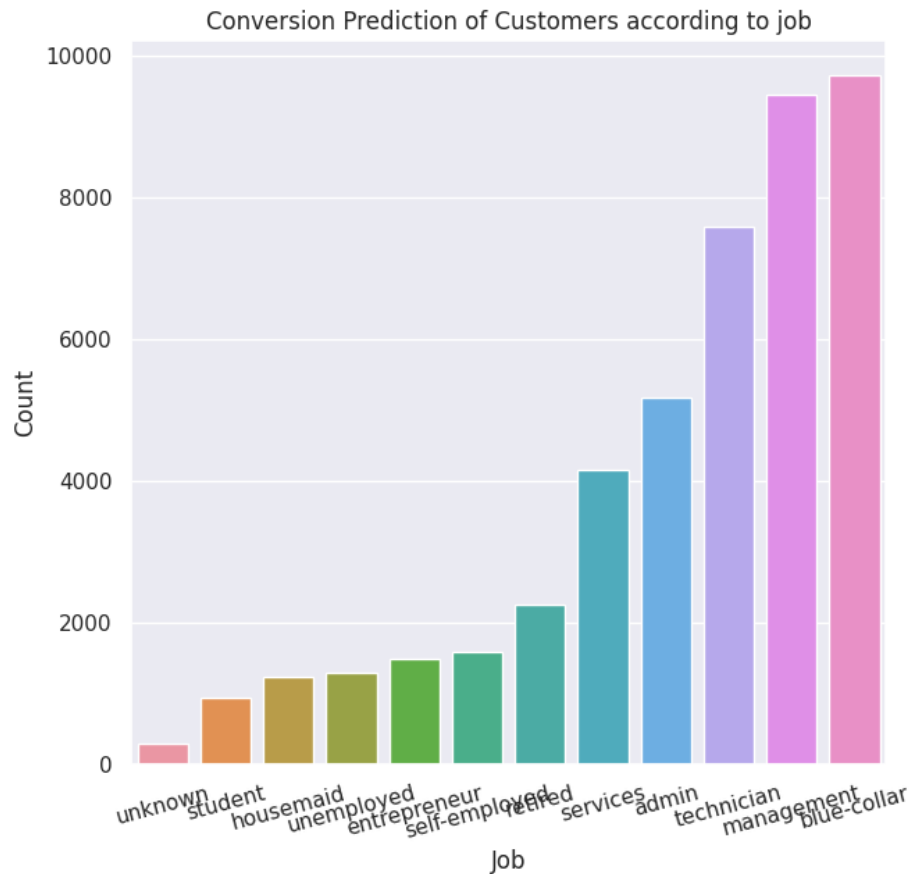
plt.subplot(436)
sns.countplot(df['y'],order=df['y'].value_counts().index[::-1])
plt.xticks(rotation=15)
plt.xlabel('Customer Conversion')
plt.ylabel('Count')
plt.title('Conversion Prediction of Customers according to Target Plot')

plt.tight_layout()
```

```

/usr/local/lib/python3.10/dist-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as
warnings.warn(
/usr/local/lib/python3.10/dist-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as
warnings.warn(
/usr/local/lib/python3.10/dist-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as
warnings.warn(
/usr/local/lib/python3.10/dist-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as
warnings.warn(
/usr/local/lib/python3.10/dist-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as
warnings.warn(
/usr/local/lib/python3.10/dist-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as
warnings.warn(
/usr/local/lib/python3.10/dist-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as
warnings.warn(

```



Conversion Prediction of Customers according to Call Type



Conversion Prediction of Customers according to day



5.2.Future Vs Target Plot

```
df.columns
```

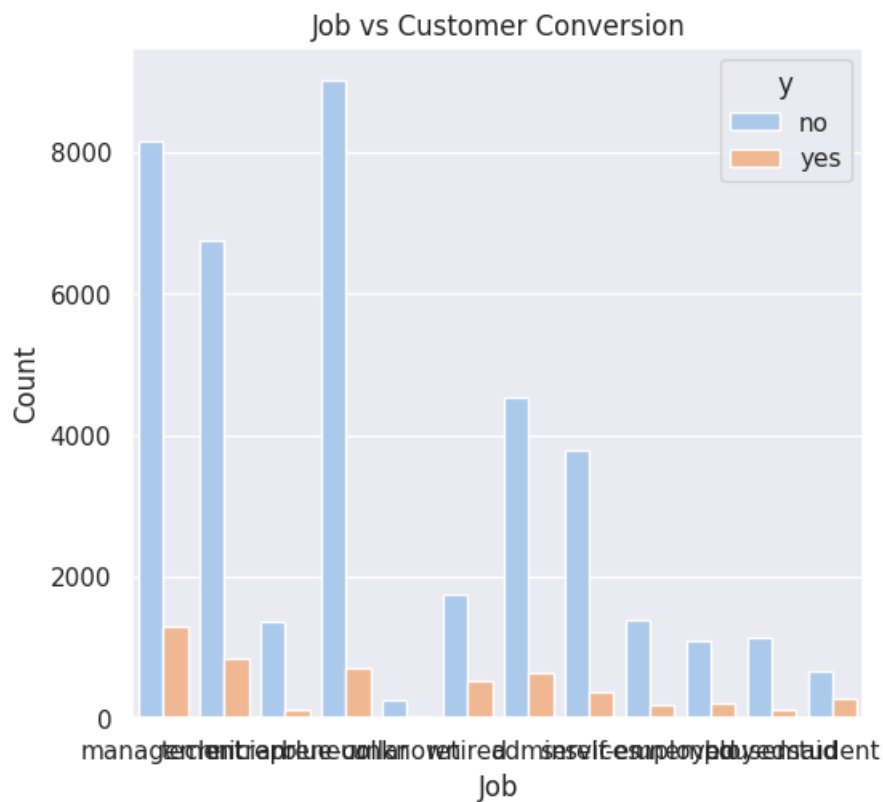
```
Index(['age', 'job', 'marital', 'education_qual', 'call_type', 'day', 'mon',
      'dur', 'num_calls', 'prev_outcome', 'y'],
      dtype='object')
```

```
plt
```

```

plt.figure(figsize=(20, 25))
plt.subplot(435)
sns.countplot(data=df, x='job', hue='y')
plt.xlabel('Job')
plt.ylabel('Count')
plt.title('Job vs Customer Conversion')
plt.show()

```



```
count_y_by_job = df.groupby('job')['y'].count().reset_index()
ordered_job_by_count = count_y_by_job.sort_values('y')['job']

plt.figure(figsize=(20, 25))
plt.subplot(431)
sns.countplot(data=df, x='job', hue='y', order=ordered_job_by_count)
plt.xlabel('Job')
plt.ylabel('Count')
plt.title('Job vs Customer Conversion')
plt.show()
```

Job vs Customer Conversion



```
#orders for all the barplots in ascending order
```

```
#Job Variable
```

```
count_y_by_job = df.groupby('job')['y'].count().reset_index()
ordered_job_by_count = count_y_by_job.sort_values('y')['job']
```

```
#marital variable
```

```
count_y_by_marital = df.groupby('marital')['y'].count().reset_index()
ordered_marital_by_count = count_y_by_marital.sort_values('y')['marital']
```

```
#education_qual Variable
```

```
count_y_by_education_qual = df.groupby('education_qual')['y'].count().reset_index()
ordered_education_qual_by_count = count_y_by_education_qual.sort_values('y')['education_qual']
```

```
#call_type Variable
```

```
count_y_by_call_type = df.groupby('call_type')['y'].count().reset_index()
ordered_call_type_by_count = count_y_by_call_type.sort_values('y')['call_type']
```

```
#month Variable
```

```
count_y_by_month = df.groupby('mon')['y'].count().reset_index()
ordered_month_by_count = count_y_by_month.sort_values('y')['mon']
```

```
#prev_outcome Variable
```

```
count_y_by_prev_outcome = df.groupby('prev_outcome')['y'].count().reset_index()
ordered_prev_outcome_by_count = count_y_by_prev_outcome.sort_values('y')['prev_outcome']
```

```
#categorical Variable
```

```
plt.figure(figsize=(20,25))
```

```
plt.subplot(431)
sns.countplot(data=df, x='job', hue='y', order=ordered_job_by_count)
plt.xlabel('Job')
plt.ylabel('Count')
plt.title('Job vs Customer Conversion')
```

```
plt.subplot(432)
sns.countplot(data=df, x='marital', hue='y', order=ordered_marital_by_count)
plt.xlabel('Marital')
plt.ylabel('Count')
plt.title('Marital vs Customer Conversion')
```

```
plt.subplot(433)
sns.countplot(data=df, x='education_qual', hue='y', order=ordered_education_qual_by_count)
plt.xlabel('Education Qualification')
plt.ylabel('Count')
plt.title('Education Qualification vs Customer Conversion')
```


```
plt.subplot(434)
sns.countplot(data=df, x='call_type', hue='y', order=ordered_call_type_by_count)
plt.xlabel('Call Type')
plt.ylabel('Count')
plt.title('Call Type vs Customer Conversion')
```

```
plt.subplot(435)
sns.countplot(data=df, x='mon', hue='y', order=ordered_month_by_count)
plt.xlabel('Month')
plt.ylabel('Count')
plt.title('Month vs Customer Conversion')
```



```
plt.subplot(436)
sns.countplot(data=df, x='prev_outcome', hue='y', order=ordered_prev_outcome_by_count)
plt.xlabel('Prev Outcome')
plt.ylabel('Count')
plt.title('Prev Outcome vs Customer Conversion')

plt.tight_layout()
```



```

Index(['age', 'job', 'marital', 'education_qual', 'call_type', 'day', 'mon',
      'dur', 'num_calls', 'prev_outcome', 'y'],
      dtype='object')

def group(x):

    if x <18:
        status= "< 18"
    elif x<24:
        status = "19-24"
    elif x<34:
        status = "25-34"
    elif x<44:
        status = "35-44"
    elif x<54:
        status = "45-54"
    else:
        status = "55 +"

    return status

df['Age_group'] = df['age'].apply(group)

import pandas as pd
import plotly.express as px

def table(df, row, column, title, xaxis_title, yaxis_title):
    def bivariate_table(df, row, column):
        table = df.groupby([row, column]).size().reset_index()
        table['percentage'] = df.groupby([row, column]).size().groupby(level=0).apply(lambda x: 100 * x / float(x.sum))
        table.columns = [row, column, 'Counts', 'Percentage']
        table['Percentage'] = table['Percentage'].astype(str) + '%'
        return table

    table = bivariate_table(df, row, column)

    fig = px.bar(table, x=row, y='Counts', color=column, barmode='stack', text=table['Percentage'])
    fig.update_layout(title=title, xaxis_title=xaxis_title, yaxis_title=yaxis_title, width=800, height=600)

    fig.show()

    return table

table(df, 'Age_group', 'y', 'Age vs Customer Conversion', 'age', 'Percentage %')
```


<ipython-input-32-4215a39c0ed2>:7: FutureWarning: Not prepending group keys to the result index of transform-like
To preserve the previous behavior, use

```
>>> .groupby(..., group_keys=False)
```

To adopt the future behavior and silence this warning, use

```
>>> .groupby(..., group_keys=True)  
table['percentage'] = df.groupby([row, column]).size().groupby(level=0).apply(lambda x: 100 * x / float(x.sum(
```

Age vs Customer Conversion



	Age_group	y	Counts	Percentage
0	19-24	no	368	72.6%

```
def day_group(day_D):  
    if day_D <= 5:  
        return "<5 Days"  
    elif day_D <= 10:  
        return "6-10 Days"  
    elif day_D <= 15:  
        return "11-15 Days"  
    elif day_D <= 20:  
        return "16-20 Days"  
    elif day_D <= 25:  
        return "21-25 Days"  
    else:  
        return "25 + Days"  
  
df['Day_Group'] = df['day'].apply(day_group)  
df['Day_Group'] = pd.Categorical(df['Day_Group'], categories=['<5 Days', '6-10 Days', '11-15 Days', '16-20 Days', '21-25 Days', '25 + Days'])  
  
table(df, 'Day_Group', 'y', 'Day vs Customer Conversion', 'Day', 'Percentage %')
```



```
<ipython-input-32-4215a39c0ed2>:7: FutureWarning:

def dur_group(dur):
    if dur <= 700: # 700 seconds = 11 minutes 40 seconds
        return "11-23 minutes"
    elif dur <= 1400: # 1400 seconds = 23 minutes 20 seconds
        return "23-40 minutes"
    elif dur <= 2100: # 2100 seconds = 35 minutes
        return "40-60 minutes"
    elif dur <= 3600: # 3600 seconds = 60 minutes (1 hour)
        return "1-2 hours"
    else:
        return "Over 2 hours"

df['Duration_Group'] = df['dur'].apply(dur_group)
df['Duration_Group'] = pd.Categorical(df['Duration_Group'],
                                     categories=['11-23 minutes', '23-40 minutes', '40-60 minutes', '1-2 hours', 'Over 2 hours'],
                                     ordered=True)

table(df, 'Duration_Group', 'y', 'Duration of Call vs Customer Conversion', 'Duration', 'Percentage %')
```

```
<ipython-input-32-4215a39c0ed2>:7: FutureWarning:
```

Not prepending group keys to the result index of transform-like apply. In the future, the group keys will be included. To preserve the previous behavior, use

```
>>> .groupby(..., group_keys=False)
```

To adopt the future behavior and silence this warning, use

```
>>> .groupby(..., group_keys=True)
```

```
def calls_group(num_calls):
    if num_calls <= 15: # Grouping: 1-15 calls
        return "1-15 calls"
    elif num_calls <= 30: # Grouping: 16-30 calls
        return "16-30 calls"
    elif num_calls <= 45: # Grouping: 31-45 calls
        return "31-45 calls"
    elif num_calls <= 60: # Grouping: 46-60 calls
        return "46-60 calls"
    else:
        return "Over 60 calls"

df['Calls_Group'] = df['num_calls'].apply(calls_group)
df['Calls_Group'] = pd.Categorical(df['Calls_Group'],
                                  categories=['1-15 calls', '16-30 calls', '31-45 calls', '46-60 calls', 'Over 60 calls'],
                                  ordered=True)
```

—

```
table(df, 'Calls_Group', 'y', 'Number of Calls vs Customer Conversion', 'Number of Calls', 'Percentage %')
```

```
<ipython-input-32-4215a39c0ed2>:7: FutureWarning:
```


Not prepending group keys to the result index of transform-like apply. In the future, the group keys will be included. To preserve the previous behavior, use

```
>>> .groupby(..., group_keys=False)
```

To adopt the future behavior and silence this warning, use

```
>>> .groupby(..., group_keys=True)
```

Number of Calls vs Customer Conversion



5.3. Encoding Categorical Variables

```
df.columns
```

```
Index(['age', 'job', 'marital', 'education_qual', 'call_type', 'day', 'mon',
      'dur', 'num_calls', 'prev_outcome', 'y', 'Age_group', 'Day_Group',
      'Duration_Group', 'Calls_Group'],
      dtype='object')
```

```
df['prev_outcome'].value_counts()
```

```
unknown    36953
failure     4901
other       1840
success     1511
Name: prev_outcome, dtype: int64
```

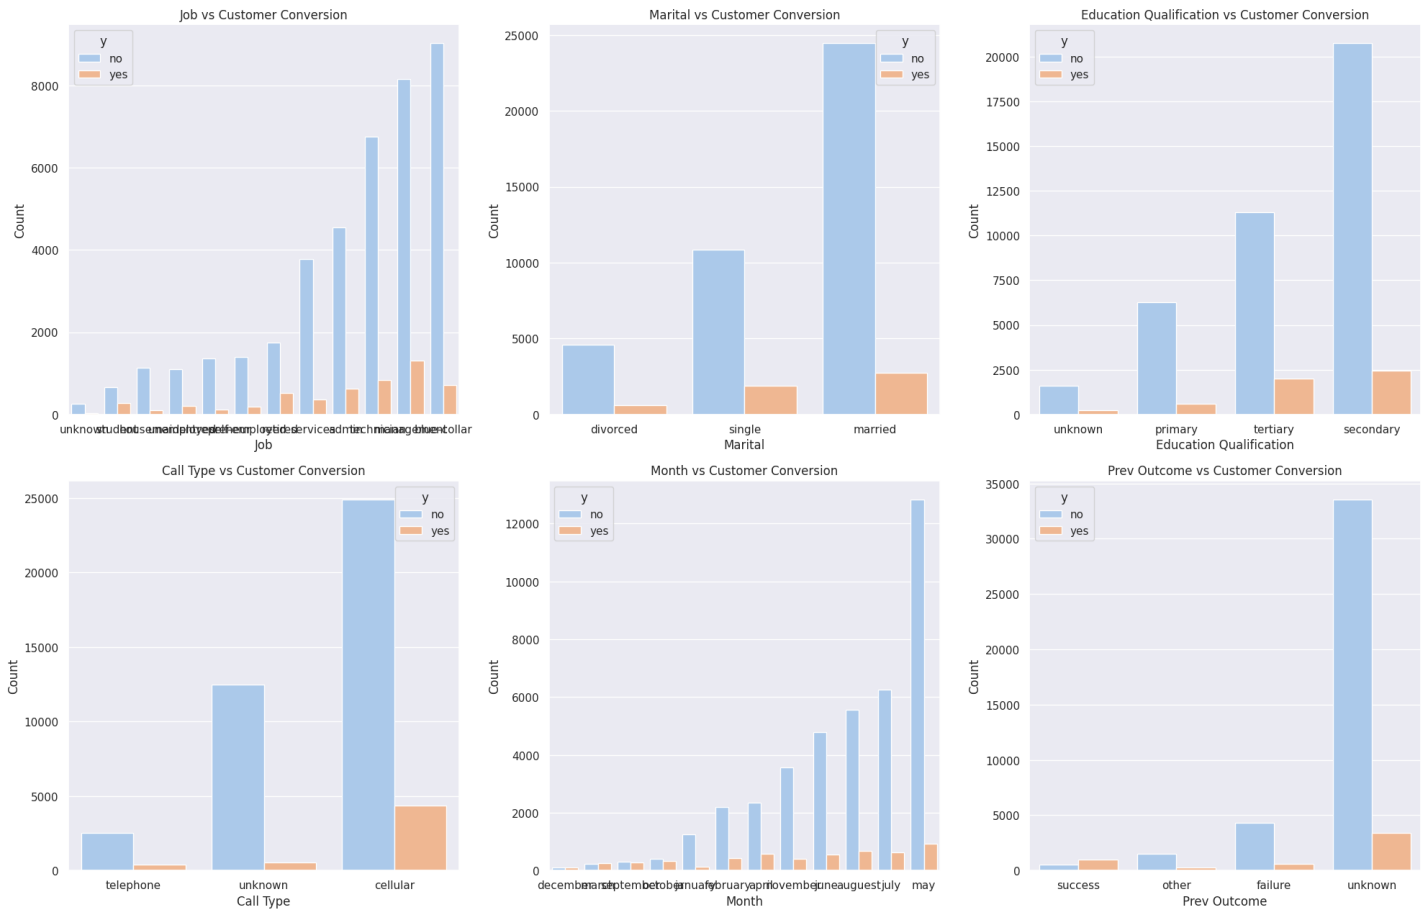
1-15 calls

16-30 calls

31-45 calls

46-60 calls

Over 60 calls



```
df['job'] = df['job'].map({'unknown':0, 'student':1, 'housemaid':2, 'unemployed':3, 'entrepreneur':4, 'self-employed':5, 'other':6})
df['marital'] = df['marital'].map({'divorced':0, 'single':1, 'married':2})
df['education_qual'] = df['education_qual'].map({'unknown':0, 'primary':1, 'tertiary':2, 'secondary':3})
df['call_type'] = df['call_type'].map({'telephone':0, 'unknown':1, 'cellular':2})
df = pd.get_dummies(df, columns = ['mon'])
##df['mon'] = df['mon'].map({'december':0, 'march':1, 'september':2, 'october':3, 'january':4, 'february':5, 'april':6, 'november':7})
df = pd.get_dummies(df, columns = ['prev_outcome'])
##df['prev_outcome'] = df['prev_outcome'].map({'success':0, 'other':1, 'failure':2, 'unknown':3})
df['y'] = df['y'].map({'yes':0, 'no':1})
```

```
df.head()
```

	age	job	marital	education_qual	call_type	day	dur	num_calls	y	Age_group	Day_Group	Duration_Group	Cal
0	58	10	2	2	1	5	261	1	1	55 +	<5 Days	11-23 minutes	
1	44	9	1	3	1	5	151	1	1	45-54	<5 Days	11-23 minutes	
2	33	4	2	3	1	5	76	1	1	25-34	<5 Days	11-23 minutes	
3	47	11	2	0	1	5	92	1	1	45-54	<5 Days	11-23 minutes	
4	33	0	1	0	1	5	198	1	1	25-34	<5 Days	11-23 minutes	

6. Splitting the data into input data and output data


```
df.columns
```

```
Index(['age', 'job', 'marital', 'education_qual', 'call_type', 'day', 'dur',
      'num_calls', 'y', 'Age_group', 'Day_Group', 'Duration_Group',
      'Calls_Group', 'mon_april', 'mon_auguest', 'mon_december',
      'mon_february', 'mon_january', 'mon_july', 'mon_june', 'mon_march',
      'mon_may', 'mon_november', 'mon_october', 'mon_september',
      'prev_outcome_failure', 'prev_outcome_other', 'prev_outcome_success',
      'prev_outcome_unknown'],
      dtype='object')
```

```
del df['Age_group']
del df['Day_Group']
del df['Duration_Group']
del df['Calls_Group']
```

```
df.columns
```

```
Index(['age', 'job', 'marital', 'education_qual', 'call_type', 'day', 'dur',
      'num_calls', 'y', 'mon_april', 'mon_auguest', 'mon_december',
      'mon_february', 'mon_january', 'mon_july', 'mon_june', 'mon_march',
      'mon_may', 'mon_november', 'mon_october', 'mon_september',
      'prev_outcome_failure', 'prev_outcome_other', 'prev_outcome_success',
      'prev_outcome_unknown'],
      dtype='object')
```

```
df.head()
```

	age	job	marital	education_qual	call_type	day	dur	num_calls	y	mon_april	mon_auguest	mon_december	mor
0	58	10	2	2	1	5	261	1	1	0	0	0	
1	44	9	1	3	1	5	151	1	1	0	0	0	
2	33	4	2	3	1	5	76	1	1	0	0	0	
3	47	11	2	0	1	5	92	1	1	0	0	0	
4	33	0	1	0	1	5	198	1	1	0	0	0	

```
def split (dataframe):
    x = dataframe.drop('y',axis=1)
    y = dataframe['y']
    return x,y
```

```
x,y = split(df)
```

```
x
```

	age	job	marital	education_qual	call_type	day	dur	num_calls	mon_april	mon_auguest	mon_december	
0	58	10	2	2	1	5	261	1	0	0	0	
1	44	9	1	3	1	5	151	1	0	0	0	
2	33	4	2	3	1	5	76	1	0	0	0	
3	47	11	2	0	1	5	92	1	0	0	0	
4	33	0	1	0	1	5	198	1	0	0	0	
...	

y

0	1
1	1
2	1
3	1
4	1
...	...
45206	0
45207	0
45208	0
45209	1
45210	1

Name: y, Length: 45205, dtype: int64

```
#splitting the data into training and testing sets with the ratio of 8:2
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(x,y,test_size=0.2,random_state=70)

print(X_train.shape,X_test.shape,y_train.shape,y_test.shape)
```

(36164, 24) (9041, 24) (36164,) (9041,)

X_train

	age	job	marital	education_qual	call_type	day	dur	num_calls	mon_april	mon_auguest	mon_december	
2955	40	9	1	0	1	14	1028	2	0	0	0	
18005	49	4	0	0	2	30	988	2	0	0	0	
36166	45	11	2	0	2	11	306	1	0	0	0	
3657	55	9	2	3	1	16	248	3	0	0	0	
11307	38	5	2	3	1	18	313	2	0	0	0	
...	
21569	57	2	2	1	2	19	111	4	0	1	0	
25922	54	3	2	2	2	19	394	2	0	0	0	
44830	81	6	2	1	0	17	231	1	0	0	0	
21624	56	9	2	3	2	19	35	4	0	1	0	
23892	33	7	2	3	2	29	155	2	0	1	0	

36164 rows × 24 columns

7.Building Machine Learning Model

```

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
import xgboost as xgb
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB

from sklearn.metrics import accuracy_score, classification_report
#Logistic Regression
#Random Forest
#Gradient Boosting (XGBoost)
#K-Nearest Neighbors (KNN)
#Naive Bayes

```

▼ 7.1 Logistic Regression

```

logistic_reg = LogisticRegression()
logistic_reg.fit(X_train, y_train)

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning:

lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

▼ LogisticRegression
LogisticRegression()

```

```

predictions_lr = logistic_reg.predict(X_test)

```

```

accuracy_lr = accuracy_score(y_test, predictions_lr)
print(f"Logistic Regression Accuracy: {accuracy_lr}")
print("Logistic Regression Classification Report:")
print(classification_report(y_test, predictions_lr))

```

```

Logistic Regression Accuracy: 0.8976883088153965
Logistic Regression Classification Report:

```

	precision	recall	f1-score	support
0	0.62	0.32	0.43	1056
1	0.92	0.97	0.94	7985
accuracy			0.90	9041
macro avg	0.77	0.65	0.68	9041
weighted avg	0.88	0.90	0.88	9041

▼ 7.2 Random Forest:

```

# Train the Random Forest model
random_forest = RandomForestClassifier()
random_forest.fit(X_train, y_train)

```

```

RandomForestClassifier
RandomForestClassifier()

```

```
predictions_rf = random_forest.predict(X_test)
```

```

accuracy_rf = accuracy_score(y_test, predictions_rf)
print(f"\nRandom Forest Accuracy: {accuracy_rf}")
print("Random Forest Classification Report:")
print(classification_report(y_test, predictions_rf))

```

```

Random Forest Accuracy: 0.9029974560336246
Random Forest Classification Report:

```

	precision	recall	f1-score	support
0	0.62	0.45	0.52	1056
1	0.93	0.96	0.95	7985
accuracy			0.90	9041
macro avg	0.77	0.71	0.73	9041
weighted avg	0.89	0.90	0.90	9041

▼ 7.3 Gradient Boosting (XGBoost):

```

# Train the XGBoost model
xgboost = xgb.XGBClassifier()
xgboost.fit(X_train, y_train)

```

```

XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=None, max_leaves=None,
               min_child_weight=None, missing=nan, monotone_constraints=None,
               multi_strategy=None, n_estimators=None, n_jobs=None,
               num_parallel_tree=None, random_state=None, ...)

```

```
predictions_xgb = xgboost.predict(X_test)
```

```

accuracy_xgb = accuracy_score(y_test, predictions_xgb)
print(f"\nXGBoost Accuracy: {accuracy_xgb}")
print("XGBoost Classification Report:")
print(classification_report(y_test, predictions_xgb))

```

```

XGBoost Accuracy: 0.9058732441101648
XGBoost Classification Report:

```

	precision	recall	f1-score	support
0	0.62	0.49	0.55	1056
1	0.93	0.96	0.95	7985
accuracy			0.91	9041
macro avg	0.78	0.73	0.75	9041

weighted avg	0.90	0.91	0.90	9041
--------------	------	------	------	------

▼ 7.4 K-Nearest Neighbors (KNN):

```
# Train the KNN model
knn = KNeighborsClassifier()
knn.fit(X_train, y_train)
```

```
▼ KNeighborsClassifier
KNeighborsClassifier()
```

```
predictions_knn = knn.predict(X_test)
```

```
accuracy_knn = accuracy_score(y_test, predictions_knn)
print(f"\nK-Nearest Neighbors Accuracy: {accuracy_knn}")
print("K-Nearest Neighbors Classification Report:")
print(classification_report(y_test, predictions_knn))
```

```
K-Nearest Neighbors Accuracy: 0.8822032960955647
K-Nearest Neighbors Classification Report:
```

	precision	recall	f1-score	support
0	0.49	0.23	0.31	1056
1	0.90	0.97	0.94	7985
accuracy			0.88	9041
macro avg	0.70	0.60	0.62	9041
weighted avg	0.86	0.88	0.86	9041

▼ 7.5 Naive Bayes

```
# Train the Naive Bayes model
naive_bayes = GaussianNB()
naive_bayes.fit(X_train, y_train)
```

```
▼ GaussianNB
GaussianNB()
```

```
predictions_nb = naive_bayes.predict(X_test)
```

```
accuracy_nb = accuracy_score(y_test, predictions_nb)
print(f"\nNaive Bayes Accuracy: {accuracy_nb}")
print("Naive Bayes Classification Report:")
print(classification_report(y_test, predictions_nb))
```

```
Naive Bayes Accuracy: 0.8816502599269992
Naive Bayes Classification Report:
```

	precision	recall	f1-score	support
0	0.49	0.41	0.45	1056
1	0.92	0.94	0.93	7985

accuracy			0.88	9041
macro avg	0.71	0.68	0.69	9041
weighted avg	0.87	0.88	0.88	9041

▼ 7.6 Cross-validation scores

```
from sklearn.model_selection import cross_val_score

models = [logistic_reg, random_forest, xgboost, knn, naive_bayes]
model_names = ['Logistic Regression', 'Random Forest', 'XGBoost', 'K-Nearest Neighbors', 'Naive Bayes']

for model, name in zip(models, model_names):
    cv_scores = cross_val_score(model, X_train, y_train, cv=10)
    print(f"{name} - Cross Validation Score: {cv_scores.mean()} (±{cv_scores.std()})")
```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning:

lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning:

lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning:

lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning:

lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning:

lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning:
```

```
lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

8.F1 Score for Machine-Learning Models

```
from sklearn.metrics import f1_score
```

```
models = [logistic_reg, random_forest, xgboost, knn, naive_bayes]
```

```
model_names = ['Logistic Regression Model', 'Random Forest Model', 'XGBoost Model', 'K-Nearest Neighbors Model', 'Naive Bayes Model']
```

```
for model, name in zip(models, model_names):
```

```
    model.fit(X_train, y_train)
```

```
    predictions = model.predict(X_test)
```

```
    f1 = f1_score(y_test, predictions)
```

```
    print(f"{name} - F1 Score: {f1}")
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning:
```

```
lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
Logistic Regression Model - F1 Score: 0.9438406896970434
```

```
Random Forest Model - F1 Score: 0.9463780592793014
```

```
XGBoost Model - F1 Score: 0.9474334424609302
```

```
K-Nearest Neighbors Model - F1 Score: 0.9355989599080848
```

```
Naive Bayes Model - F1 Score: 0.9337050805452293
```

9.Suggestion

```
xgboost.feature_importances_
```

```
array([0.01484445, 0.01221474, 0.01210118, 0.01092021, 0.05915434,
        0.01778406, 0.05698902, 0.01287511, 0.02980603, 0.02077172,
        0.02271581, 0.03245021, 0.02717924, 0.03175417, 0.06404951,
        0.09638961, 0.0454943 , 0.03021196, 0.06405973, 0.05480836,
        0.01264754, 0.01519127, 0.22764911, 0.02793836], dtype=float32)
```

```
df.columns
```

```
Index(['age', 'job', 'marital', 'education_qual', 'call_type', 'day', 'dur',
        'num_calls', 'y', 'mon_april', 'mon_august', 'mon_december',
        'mon_february', 'mon_january', 'mon_july', 'mon_june', 'mon_march',
        'mon_may', 'mon_november', 'mon_october', 'mon_september',
        'prev_outcome_failure', 'prev_outcome_other', 'prev_outcome_success',
        'prev_outcome_unknown'],
      dtype='object')
```

```
sorted_idx = xgboost.feature_importances_.argsort()
plt.figure(figsize=(10,5))
plt.barh(df.columns[sorted_idx], xgboost.feature_importances_[sorted_idx])
plt.xlabel("Extreme Gradient Boosting Feature Importance")
plt.title("Feature Importance")
plt.show()
```

